

W15_LSTMConsumptionSimpleV13

January 12, 2021

```
[1]: version = "13"

stationary = False
scaling = True
resetting = True

hardcode = False
train_for = 1501 #2001
learningrate = 1e-3 #1e-3
```

1 TO DO:

- Revisit train/validate/split
- Create dummy variables of the hour and day
- Add features
 - isPeak
 - Mean of previous Day
 - Mean of previous Week

2 Simpler form of the LSTM (from Week 14)

Jefry el Bhwash 16095065 (Created v01 and onward)

Niels van Drunen (joined from 03m left@05 for the production variant)

Levy Duivenvoorden (joined @03, left @05)

Using: - DL Lectures: 6.3 RNN Stationary - W14_Simple_LSTM_Consumption -
W13_LSTM_Start_shaping - Help session of mr Vuurens

Data: - House 28

Versions:

| nr | Date | Changes |
|----|-----------|---|
| 01 | 10/12/'20 | First draft, does not work |
| 02 | 11/12/'20 | Added validation (valid set, no actual tests yet) and visualization |
| 03 | 14/12/'20 | Added this notebook to the server, added Levy and Niels, to try and understand LSTM's better as a group |
| 04 | 14/12/'20 | Added Validation and visualization during training |
| 05 | 15/12/'20 | Quick look at the test set to see if something weird is going on |
| 06 | 15/12/'20 | Removed Test, because that's what validation is for. Changed print while training. Added visualizations on the bottom and made the usage of stationary variable on the top of this notebook |
| 07 | Skipped | NvD added his own beautiful version of 07 to the server |
| 08 | 15/12/'20 | Scaled the dataframes going in |
| 09 | 16/12/'20 | Changed where Scaling happens |
| 10 | 16/12/'20 | Changed dim3 function |
| 11 | 17/12/'20 | Added multiple features, dataloader |
| 12 | 18/12/'20 | Inverse scaling fixed, added learningrate scheduler + resetting, better training visualization |

3 Initialization

```
[2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from matplotlib import gridspec

from IPython.display import display, HTML
```

```
import time
```

```
[3]: import random
      #Neural Network imports
      import torch
      import torch.nn as nn
      import torch.optim as optim
      from torch.utils.data import DataLoader, TensorDataset

      from sklearn.preprocessing import StandardScaler
      from sklearn.metrics import mean_squared_error
      from sklearn.metrics import r2_score
```

```
[4]: #cuda imports
      ngpu = torch.cuda.device_count() # number of available gpus
      device = torch.device("cuda:4") if (torch.cuda.is_available() and ngpu > 0)
      ↪else "cpu" #cuda:0 for gpu 0, cuda:4 for gpu 5

      #Random Seed
      random.seed(1337)
      torch.manual_seed(1337)

      def det(tensor):
          return tensor.detach().cpu().numpy()
```

4 Data Preparation

```
[5]: df = pd.read_pickle('consumptionOf28').drop(['production'], axis=1)
      #adding hour feature
      df['hour'] = df.index.hour
      df['weekday'] = df.index.dayofweek
      #df['day'] = df.index.day

      #AUTOMATE THIS
      cols = ['hour', 'weekday', 'consumption']
      df = df[cols]
      display(df.head(2))
```


| | hour | weekday | consumption |
|---------------------|------|---------|-------------|
| 2018-12-31 23:00:00 | 23 | 0 | 0.573 |
| 2019-01-01 00:00:00 | 0 | 1 | 0.435 |

4.1 Data Split

```
[6]: # Split
trdf = df.loc['2019-01':'2019-08']
vadf = df.loc['2019-09':'2019-11']
tedf = df.loc['2019-12':]
```

```
[7]: plt.plot(trdf.index, trdf.consumption, '.-', alpha=0.5)
plt.plot(vadf.index, vadf.consumption, '.-', alpha=0.5)
plt.plot(tedf.index, tedf.consumption, '.-', alpha=0.5)
```

```
[7]: [<matplotlib.lines.Line2D at 0x7f1b1a6e7358>]
```



4.2 Scaling

```
[8]: dftr = trdf.copy()
dfva = vadf.copy()
dfte = tedf.copy()
if scaling:
    print('Scaling: On')

    scaler_X = StandardScaler()
    scaler_y = StandardScaler()
```

```

scaler_X.fit(dftr.iloc[:, :-1])
scaler_y.fit(dftr.iloc[:, -1:])

#train
dftr.iloc[:, :-1] = scaler_X.transform(dftr.iloc[:, :-1])
dftr.iloc[:, -1:] = scaler_y.transform(dftr.iloc[:, -1:])
print('Train')
display(dftr.head())

#Valid
dfva.iloc[:, :-1] = scaler_X.transform(dfva.iloc[:, :-1])
dfva.iloc[:, -1:] = scaler_y.transform(dfva.iloc[:, -1:])
print(f'\nValidation')
display(dfva.head())

#Test
dfte.iloc[:, :-1] = scaler_X.transform(dfte.iloc[:, :-1])
dfte.iloc[:, -1:] = scaler_y.transform(dfte.iloc[:, -1:])
print(f'\nTest')
display(dfte.head())

```

Scaling: On

Train

| | hour | weekday | consumption |
|---------------------|-----------|-----------|-------------|
| 2019-01-01 00:00:00 | -1.661325 | -1.005184 | -0.305184 |
| 2019-01-01 01:00:00 | -1.516862 | -1.005184 | -0.739154 |
| 2019-01-01 02:00:00 | -1.372399 | -1.005184 | -0.999535 |
| 2019-01-01 03:00:00 | -1.227936 | -1.005184 | 2.293371 |
| 2019-01-01 04:00:00 | -1.083473 | -1.005184 | -0.473512 |

Validation

| | hour | weekday | consumption |
|---------------------|-----------|----------|-------------|
| 2019-09-01 00:00:00 | -1.661325 | 1.507776 | -1.002165 |
| 2019-09-01 01:00:00 | -1.516862 | 1.507776 | -0.999535 |
| 2019-09-01 02:00:00 | -1.372399 | 1.507776 | -0.802276 |
| 2019-09-01 03:00:00 | -1.227936 | 1.507776 | -0.313075 |
| 2019-09-01 04:00:00 | -1.083473 | 1.507776 | 1.675294 |

Test

| | hour | weekday | consumption |
|---------------------|-----------|----------|-------------|
| 2019-12-01 00:00:00 | -1.661325 | 1.507776 | 0.802095 |
| 2019-12-01 01:00:00 | -1.516862 | 1.507776 | 0.441769 |
| 2019-12-01 02:00:00 | -1.372399 | 1.507776 | 0.483851 |
| 2019-12-01 03:00:00 | -1.227936 | 1.507776 | 0.660069 |
| 2019-12-01 04:00:00 | -1.083473 | 1.507776 | 1.186092 |

4.3 3 dimensional

batch | sequence | features

```
[9]: def dim3(dft, window=7, gap=24):  
    #Get time shifted values and apply a moving window  
    X = np.concatenate([ dft[i:i+window].to_numpy().reshape(1, window, dft.  
→shape[1]) for i in range(len(dft)-window-gap) ], axis=0)  
  
    #Get the target value (which is the next one in the sequence)  
    y = dft.to_numpy()[window + gap:, -1]  
    print(f"X_shape: {X.shape}")  
    print(f"y_shape: {y.shape}")  
    return X.astype(np.float32), y.astype(np.float32)
```

```
[10]: wsize = 168  
      gsize = 24  
  
      print("Train")  
      train_X, train_y = dim3(dftr, wsize, gsize)  
      print("\nValid")  
      valid_X, valid_y = dim3(dfva, wsize, gsize)  
      print("\nTest")  
      test_X, test_y = dim3(dfte, wsize, gsize)
```

Train

X_shape: (5640, 168, 3)
y_shape: (5640,)

Valid

X_shape: (1992, 168, 3)
y_shape: (1992,)

Test

X_shape: (551, 168, 3)
y_shape: (551,)

```
[11]: display(train_X[0][:10])  
      print("...")  
      display(train_y[0])
```

```
array([[ -1.6613247 , -1.005184  , -0.30518422],  
       [ -1.5168618 , -1.005184  , -0.7391535 ],  
       [ -1.3723987 , -1.005184  , -0.9995351 ],  
       [ -1.2279357 , -1.005184  ,  2.2933714 ],  
       [ -1.0834727 , -1.005184  , -0.4735117 ],  
       [ -0.93900967, -1.005184  ,  0.64954823],  
       [ -0.7945466 , -1.005184  ,  0.49700147],
```

```

[-0.6500836 , -1.005184 ,  1.2781461 ],
[-0.5056206 , -1.005184 , -0.71285236],
[-0.36115757, -1.005184 , -1.0626578 ]], dtype=float32)

...

-0.85487866

```

5 Tensors and Dataloader

Dataloader for train

Tensor for valid and test

```

[12]: train_ds = TensorDataset(torch.tensor(train_X), torch.tensor(train_y))
      train_dl = DataLoader(train_ds, batch_size=64, num_workers=3)

      valid_X_t = torch.from_numpy(np.array(valid_X)).to(device)
      valid_y_t = torch.from_numpy(np.array(valid_y)).to(device)

      test_X_t = torch.from_numpy(np.array(test_X)).to(device)
      test_y_t = torch.from_numpy(np.array(test_y)).to(device)

```

6 LSTM Class

```

[13]: class lstm(nn.Module):
      def __init__(self, feature_size=1, hidden_state_size = 100):
          super().__init__()
          self.hidden_state_size = hidden_state_size
          self.lstm1 = nn.LSTM(feature_size, self.hidden_state_size,
      ↪batch_first=True)
          self.linear2 = nn.Linear(self.hidden_state_size, 1)

      def forward(self, X): #tensor X
          h, _ = self.lstm1( X )          # h shaped (batch, sequence,
      ↪hidden_layer)
          h = h[:, -1, :]                # only need the output for the last
      ↪sequence

          y = self.linear2(h)            # make a prediction
          if stationary:
              y = y + X[:, -1, -1:]      # make the output stationary
          return y.view(-1)              # like always

```

6.0.1 LSTM Object

```
[14]: model = lstm(df.shape[1]).to(device)
```

7 LSTM training

```
[15]: stime = time.time()
show_every = 100
reset_scheduler_after_n_epochs = 100
if hardcode:
    train_for = 2001
    learningrate = 1e-3

#visualization parameters
alijst = []; blijst = []; lrlijst = [];

#Training parameters:
optimizer = optim.AdamW(model.parameters(), lr=learningrate)
scheduler = torch.optim.lr_scheduler.OneCycleLR(optimizer, max_lr=learningrate,
    ↪steps_per_epoch=len(train_dl), epochs=train_for)
criterion = nn.SmoothL1Loss()

ttime = time.time()
#Learning loop:
for i in (range(train_for)):
    """
    Training
    """
    model.train()

    for X, y in train_dl:
        X, y = X.to(device), y.to(device)

        #train LSTM and reshape output:
        optimizer.zero_grad()
        output = model(X)

        #bereken de loss over de output en update de parameters:
        loss = criterion(output, y)
        lossT = mean_squared_error(det(output), det(y))
        loss.backward()
        optimizer.step()
        scheduler.step()
    """
    Evaluation
```



```

"""
model.eval()
optimizer.zero_grad()

dataV = valid_X_t;
targetV = valid_y_t.view(-1);

outputV = model(dataV)

#bereken de loss over de output en update de parameters:
lossV = mean_squared_error(det(outputV), det(targetV))

alist.append(lossT)
blijst.append(lossV)
lrlijst.append(scheduler.get_last_lr())

#Plotting the prediction on the validation set and plotting the learningrate
if i%show_every == 0:
    vtime = time.time()
    fig = plt.figure(figsize=(13,5))
    spec = gridspec.GridSpec(ncols=2, nrows=1,width_ratios=[4, 1])

    display(HTML(f'<h3 style="text-align:center"> {i} </h3>'))
    fig.suptitle(f"Consumption LSTM in {round(vtime-ttime, 2)} seconds")

    ax0 = fig.add_subplot(spec[0],title=f'Epoch: {i} of {train_for} ||  

↳LSTMVersion: {version}\nPrediction LSTM model on the Validation set',
                        xlabel='datapoint [hr]', ylabel='Energy  

↳Consumption [kWh]', xlim=[1000,1072])
    ax0.plot(det(outputV), alpha=0.5, label="Prediction")
    ax0.plot(det(targetV), alpha=0.5, label = "Target")
    ax0.legend()

    ax1 = fig.add_subplot(spec[1], title=f'Learningrate scheduler\n',
                        xlabel= 'Epoch', ylabel='Rate' )
    ax1.plot([i for i in range(0,len(lrlijst))], lrlijst, alpha=0.5,  

↳label='learning rate', c='r')
    ax1.grid()
    ax1.legend()#loc=(1.01, 0.5))

    plt.tight_layout()
    plt.show()
    ttime = time.time()
if resetting:
    if i % reset_scheduler_after_n_epochs == 0:

```

```

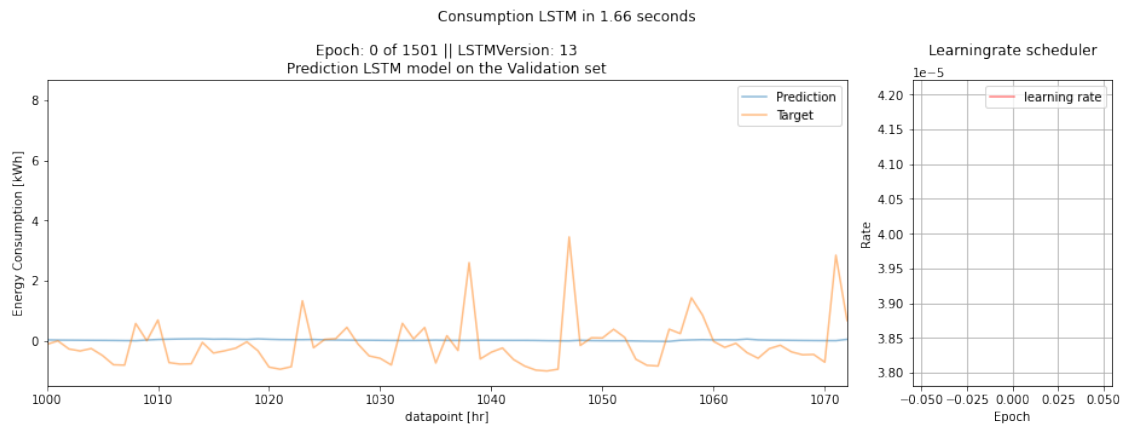
scheduler = torch.optim.lr_scheduler.OneCycleLR(optimizer,
↳max_lr=learningrate, steps_per_epoch=len(train_dl),
↳epochs=reset_scheduler_after_n_epochs)

pass


etime = time.time()

```


<IPython.core.display.HTML object>




<IPython.core.display.HTML object>




<IPython.core.display.HTML object>




<IPython.core.display.HTML object>




<IPython.core.display.HTML object>




<IPython.core.display.HTML object>




<IPython.core.display.HTML object>




<IPython.core.display.HTML object>



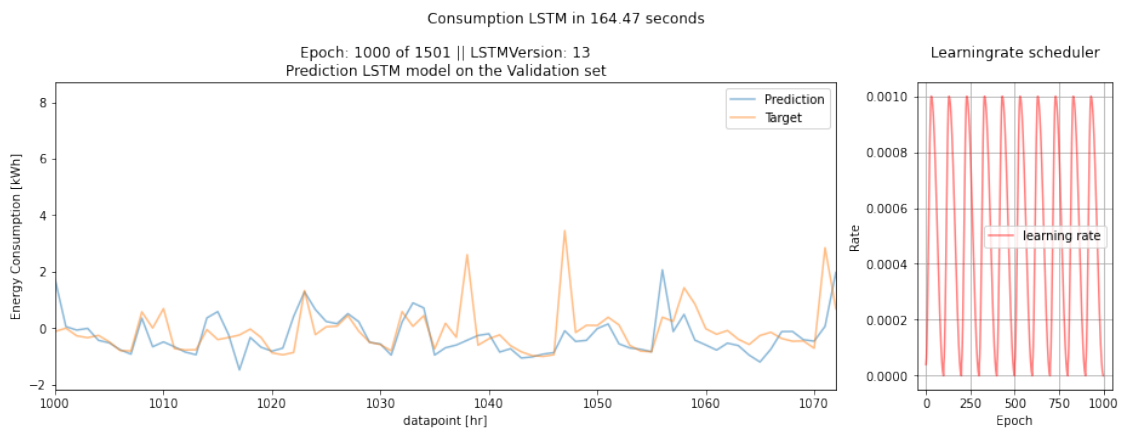
<IPython.core.display.HTML object>



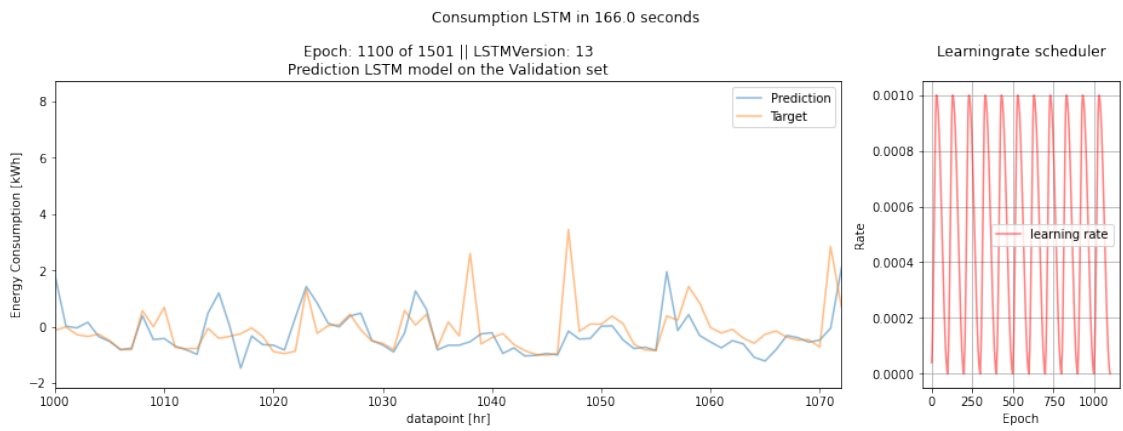
<IPython.core.display.HTML object>




<IPython.core.display.HTML object>




<IPython.core.display.HTML object>




<IPython.core.display.HTML object>




<IPython.core.display.HTML object>



<IPython.core.display.HTML object>



<IPython.core.display.HTML object>



```
[16]: print(f"Training Time for {train_for} epochs: {etime-stime}seconds")
```

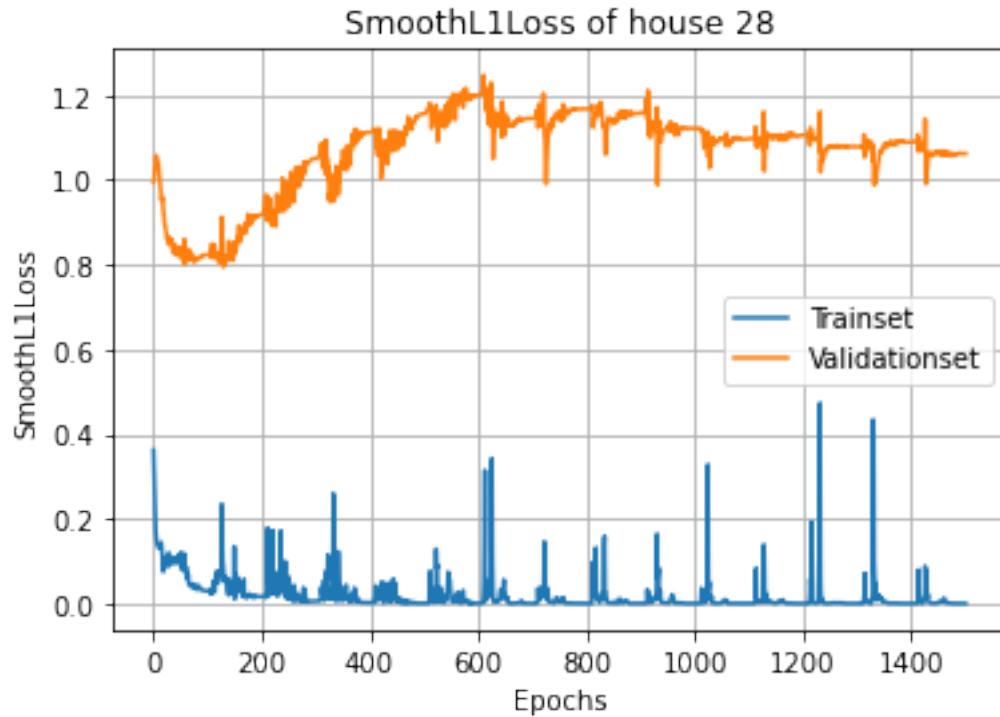
Training Time for 1501 epochs: 2456.7592170238495seconds

8 Visualizing the results

```
[17]: model.eval()

%matplotlib inline
plt.plot([i for i in range(0,len(alijst))],alijst,label="Trainset")
plt.plot([i for i in range(0,len(blijst))],blijst, label="Validationset")

plt.title(f'{str(criterion)[-2]} of house 28' )
plt.xlabel("Epochs")
plt.ylabel(f'{str(criterion)[-2]}') # Lossfunction
plt.legend()
plt.grid()
plt.show()
```

```
[18]: predictedvalue = det(outputV)
      targetvalue =det(targetV)

      if scaling:
          predictedvalue = scaler_y.inverse_transform(det(outputV))
          targetvalue = scaler_y.inverse_transform(det(targetV))

      plt.plot(predictedvalue, '.-', alpha=0.5, label="Prediction")
      plt.plot(targetvalue, '.-', alpha=0.5, label = "Target")
      plt.tight_layout()
      plt.title(f'LSTMVersion: {version} || Scaling: {"ON" if scaling else "OFF"} ||
        ↳After {train_for} Epochs || LR: {learningrate}\n\nValidation Set Target and
        ↳Prediction')
      plt.xlim([600, 672])
      #plt.ylim([0,2])
      plt.legend()
      plt.xlabel("datapoint [hr]")
      plt.ylabel("Energy [kWh]")
      plt.show()
```


```

print(f"Area under prediction: \t{predictedvalue.sum()}\nArea under Target: \t{targetvalue.sum()}\nDifference: \t\t{predictedvalue.sum()-targetvalue.sum()}\n")
print(f"Difference percentage: \t{round((predictedvalue.sum()-targetvalue.sum())/(targetvalue.sum()*100, 2)}%")

```

LSTMVersion: 13 || Scaling: ON || After 1501 Epochs || LR: 0.001

Validation Set Target and Prediction



```

Area under prediction: 1150.25244140625
Area under Target: 1196.030029296875
Difference: -45.777587890625
Difference percentage: -3.83%

```

```

[19]: y = det(valid_y_t)
      yhat = det(model(valid_X_t))


      r2 = round(r2_score(yhat, y),5)
      lossV = round(mean_squared_error(yhat, y),5)

      plt.scatter(y, yhat)
      plt.plot(plt.xlim(), plt.xlim(), ls="--", c='r', label="$y$=$\hat{y}$")

```

```
plt.title(f'R2: {r2} || MSE: {lossV} || After {train_for} Epochs || LR:␣
↳{learningrate}␣Stationary: {"ON" if stationary else "OFF"} || Scaling:␣
↳{"ON" if scaling else "OFF"} || LSTMVersion: {version}␣Criterion:␣
↳{criterion} ␣\n\nY vs $\hat{Y}$')
plt.xlabel("y [kWh]")
plt.ylabel("$\hat{Y}$ [kWh]")
plt.show()
```

R2: -0.33385 || MSE: 1.0618000030517578 || After 1501 Epochs || LR: 0.001
Stationary: OFF || Scaling: ON || LSTMVersion: 13
Criterion: SmoothL1Loss()



```
[20]: %%javascript
Jupyter.notebook.session.delete()
```

<IPython.core.display.Javascript object>